





```
[47]: param = lgbm.max
best_param = param['params']
best_param

Out[47]: {'lambda_11': 0.014904030431068804,
'max_depth': 0.024307016196213,
'max_depth': 43.31105188516426,
'min_child_samples': 82.11410102438798,
'min_data_in_leaf': 146.58612589195322,
'num_leaves': 84.1969087678835}

In [48]: for key, val in best_param.items():
if (key != 'lambda_11') & (key != 'lambda_12'):
best_param[key] = int(val)

In [49]: best_param

Out[49]: {'lambda_11': 0.014904030431068804,
'max_depth': 0.024307016196213,
'max_depth': 43,
'min_child_samples': 82,
'min_data_in_leaf': 146,
'num_leaves': 84}

In [50]: train_data = lighgbm.Dataset(X1_Train, label= y1_Train)

num_round= 4
best_lgb = lighgbm.train(best_param, train_data, num_round)
y1_Train_pred = best_lgb.predict(X1_Train)
y1_Test_pred = best_lgb.predict(X1_Test)

for i in np.arange(len(y1_Train_pred)):
if y1_Train_pred[i] >= 0.501:
y1_Train_pred[i]= 1
else:
y1_Train_pred[i]= 0

for i in np.arange(len(y1_Test_pred)):
if y1_Test_pred[i] >= 0.50:
y1_Test_pred[i]= 1
else:
y1_Test_pred[i]= 0

eval_df = evaluation('Lightgbm on Train', y1_Train, y1_Train_pred)
results = pd.concat((results, eval_df))
eval_df = evaluation('Lightgbm on Test', y1_Test, y1_Test_pred)
results = pd.concat((results, eval_df))

[LightGBM] [Warning] Find whitespaces in feature_names, replace with underlines
[LightGBM] [Warning] min_data_in_leaf is set=146, min_child_samples=82 will be ignored. Current value: min_data_in_leaf=146
[LightGBM] [Warning] Auto-choosing col-wise multi-threading, the overhead of testing was 0.011460 sec on 0s.
You can set 'force_col_wise=true' to remove the overhead.
[LightGBM] [Info] Total Bins 3951
[LightGBM] [Info] Number of data points in the train set: 13500, number of used features: 69
[LightGBM] [Info] Start training from score 0.481852
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

lightgbm on Train
precision recall f1-score support
0 1.00 0.96 0.98 6995
1 0.96 1.00 0.98 6505

accuracy 0.98 0.98 0.98 13500
macro avg 0.98 0.98 0.98 13500
weighted avg 0.98 0.98 0.98 13500

lightgbm on Test
precision recall f1-score support
0 1.00 0.97 0.98 2998
1 0.02 1.00 0.04 2

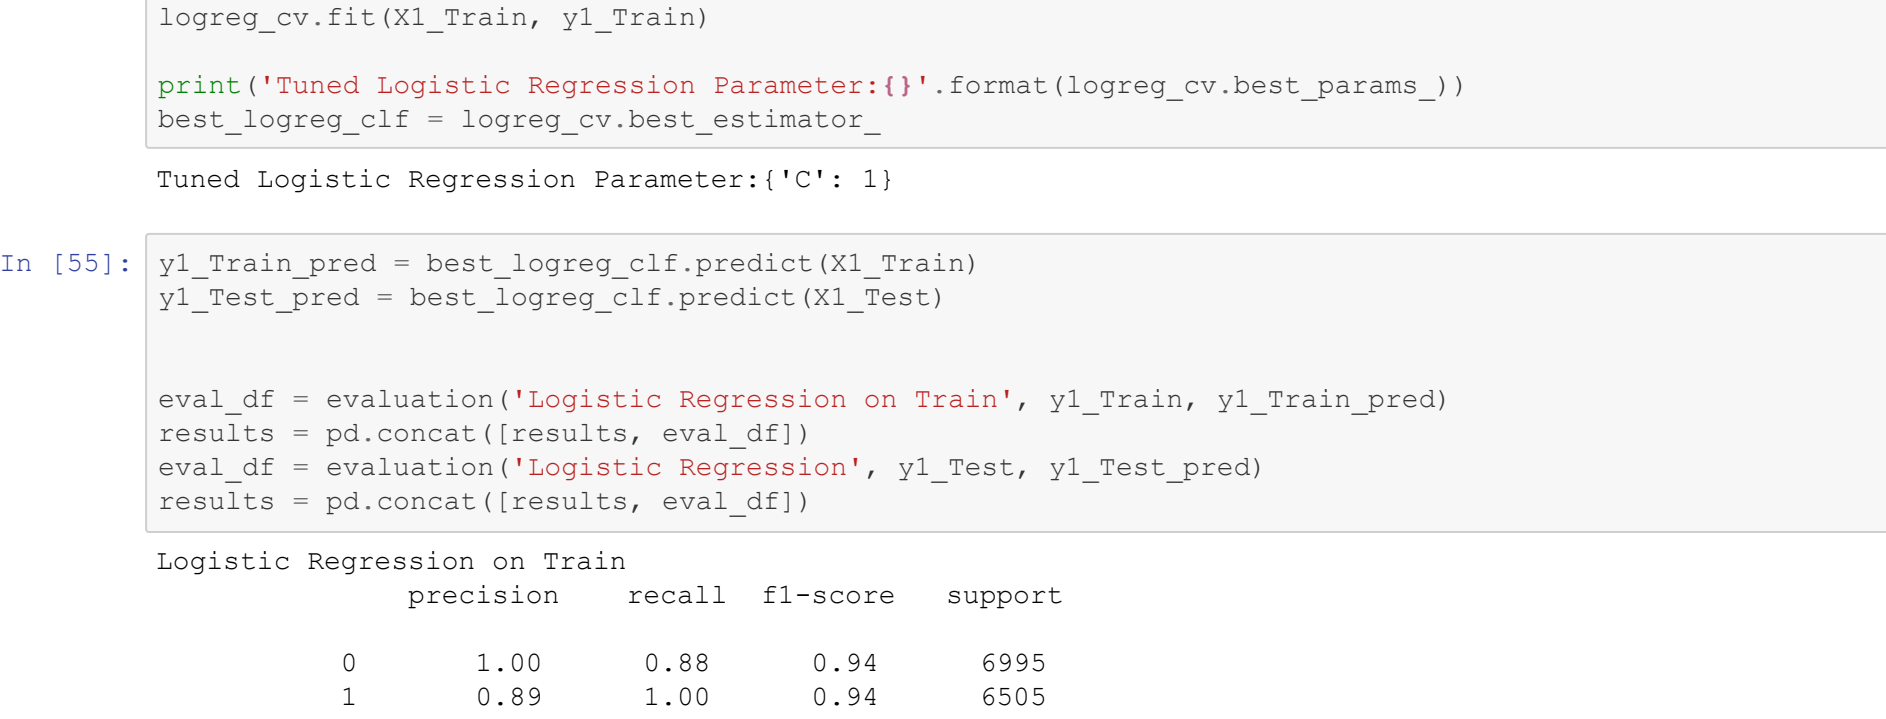
accuracy 0.51 0.98 0.51 3000
macro avg 1.00 0.97 0.98 3000
weighted avg 1.00 0.97 0.98 3000

In [51]: conf_mat(y1_Train, y1_Train_pred, 'Train', 'LGB Model')
LGB Model - Confusion Matrix on Train

Out[51]:
predicted_Non_Default predicted_Default
Non_Default 6740 255
Default 0 6505

In [52]: conf_mat(y1_Test, y1_Test_pred, 'Test', 'LGB Model')
LGB Model - Confusion Matrix on Test

Out[52]:
predicted_Non_Default predicted_Default
Non_Default 2888 100
Default 0 2
```



### Logistic regression

**Logistic Regression** is a discriminative classifier as we learn a soft boundary between/among classes.

**C** is the regularization strength (equivalent to 1/alpha from the Ridge case), and smaller values of C mean stronger regularization. As before, the regularization tries to prevent features from having terribly high weights, thus implementing a form of feature selection.

```
In [54]: from sklearn.linear_model import LogisticRegression
#the grid of parameters to search over
Cs = (0.001, 0.01, 0.1, 1, 10, 100)
param_grid = {'C': Cs}

#Instantiate a Logistic regression classifier: logreg
logreg = LogisticRegression()
# Instantiate the GridSearchCV object: logreg_cv
logreg_cv = GridSearchCV(logreg, param_grid, cv=5)
# fit it to the data
logreg_cv.fit(X1_Train, y1_Train)

print('Tuned Logistic Regression Parameter:{}'.format(logreg_cv.best_params_))
best_logreg_clf = logreg_cv.best_estimator_

Tuned Logistic Regression Parameter: {'C': 1}

In [55]: y1_Train_pred = best_logreg_clf.predict(X1_Train)
y1_Test_pred = best_logreg_clf.predict(X1_Test)

eval_df = evaluation('Logistic Regression on Train', y1_Train, y1_Train_pred)
results = pd.concat((results, eval_df))
eval_df = evaluation('Logistic Regression', y1_Test, y1_Test_pred)
results = pd.concat((results, eval_df))

Logistic Regression on Train
precision recall f1-score support
0 1.00 0.88 0.94 6995
1 0.89 1.00 0.94 6505

accuracy 0.94 0.94 0.94 13500
macro avg 0.94 0.94 0.94 13500
weighted avg 0.95 0.94 0.94 13500

Logistic Regression
precision recall f1-score support
0 1.00 0.88 0.94 2998
1 0.01 1.00 0.01 2

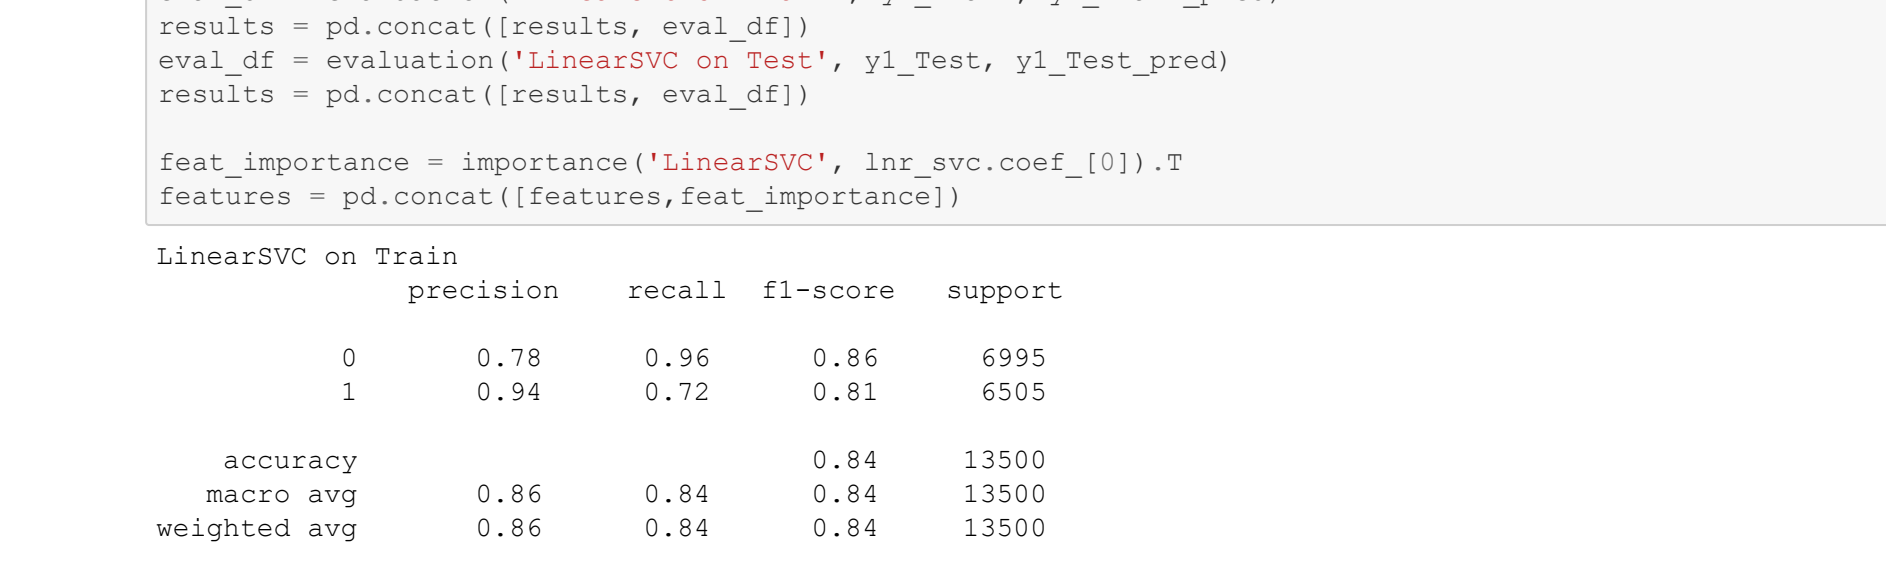
accuracy 0.50 0.94 0.47 3000
macro avg 1.00 0.88 0.94 3000
weighted avg 1.00 0.88 0.94 3000

In [56]: conf_mat(y1_Train, y1_Train_pred, 'Train', 'Logistic Regression')
Logistic Regression - Confusion Matrix on Train

Out[56]:
predicted_Non_Default predicted_Default
Non_Default 6173 822
Default 0 6505

In [57]: conf_mat(y1_Test, y1_Test_pred, 'Test', 'Logistic Regression')
Logistic Regression - Confusion Matrix on Test

Out[57]:
predicted_Non_Default predicted_Default
Non_Default 2642 358
Default 0 2
```



### LinearSVC

```
In [59]: from sklearn.svm import LinearSVC

lnr_svc = LinearSVC() #class_weight={0 : 1, 1: 49}
lnr_svc.fit(X1_Train, y1_Train)
y1_Train_pred = lnr_svc.predict(X1_Train)
y1_Test_pred = lnr_svc.predict(X1_Test)

eval_df = evaluation('LinearSVC on Train', y1_Train, y1_Train_pred)
results = pd.concat((results, eval_df))
eval_df = evaluation('LinearSVC on Test', y1_Test, y1_Test_pred)
results = pd.concat((results, eval_df))

feat_importance = importance('LinearSVC', lnr_svc.coef_[0], X1_Test.columns)
features = pd.concat((features, feat_importance))

LinearSVC on Train
precision recall f1-score support
0 0.78 0.96 0.86 6995
1 0.84 0.72 0.81 6505

accuracy 0.86 0.84 0.84 13500
macro avg 0.84 0.84 0.84 13500
weighted avg 0.86 0.84 0.84 13500

LinearSVC on Test
precision recall f1-score support
0 1.00 0.95 0.98 2998
1 0.01 1.00 0.03 2

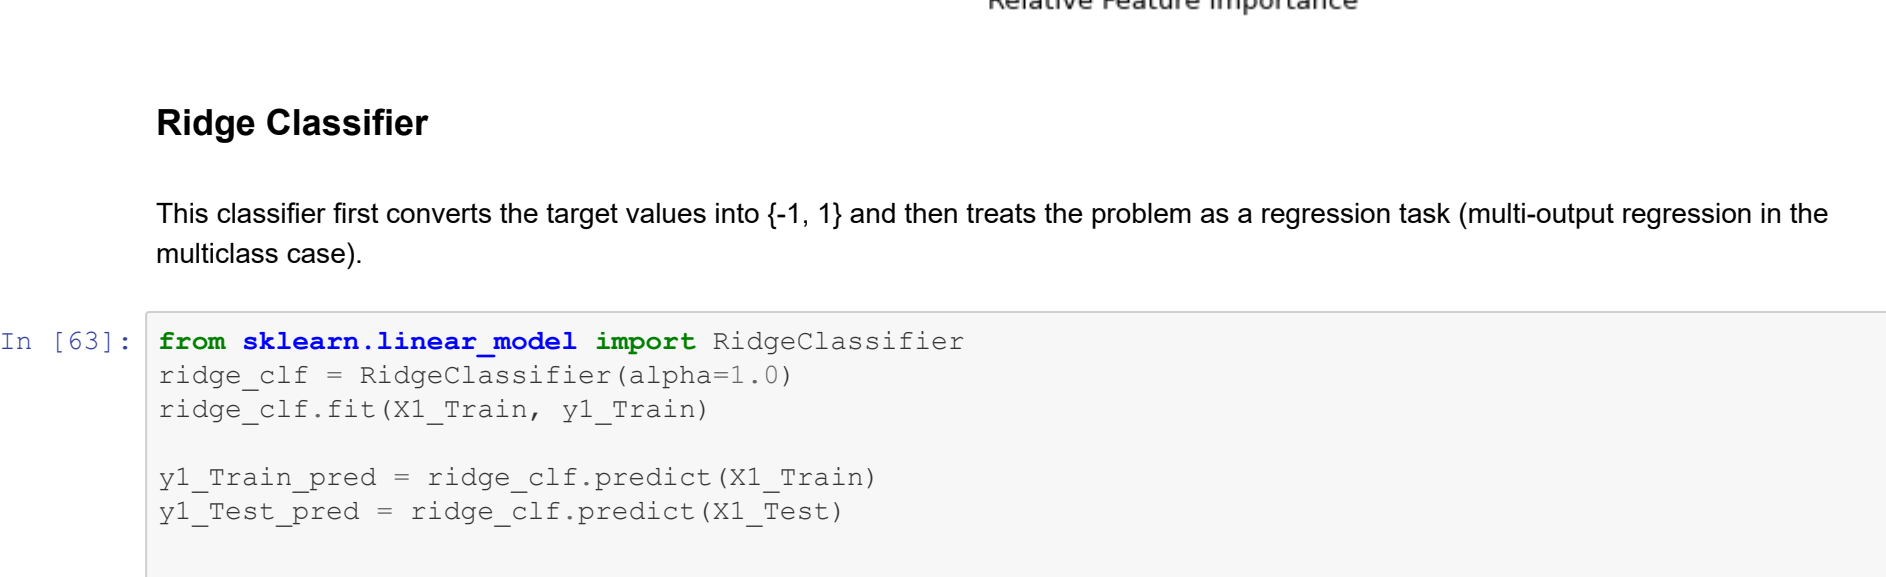
accuracy 0.51 0.98 0.50 3000
macro avg 1.00 0.95 0.98 3000
weighted avg 1.00 0.95 0.98 3000

In [60]: conf_mat(y1_Train, y1_Train_pred, 'Train', 'Linear SVC')
Linear SVC - Confusion Matrix on Train

Out[60]:
predicted_Non_Default predicted_Default
Non_Default 6714 281
Default 1651 4654

In [61]: conf_mat(y1_Test, y1_Test_pred, 'Test', 'Linear SVC')
Linear SVC - Confusion Matrix on Test

Out[61]:
predicted_Non_Default predicted_Default
Non_Default 2859 139
Default 0 2
```



### Ridge Classifier

This classifier first converts the target values into {-1, 1} and then treats the problem as a regression task (multi-output regression in the multiclass case).

```
In [63]: from sklearn.linear_model import RidgeClassifier
ridge_clf = RidgeClassifier(alpha=1.0)
ridge_clf.fit(X1_Train, y1_Train)

y1_Train_pred = ridge_clf.predict(X1_Train)
y1_Test_pred = ridge_clf.predict(X1_Test)

eval_df = evaluation('Ridge Classifier on Train', y1_Train, y1_Train_pred)
results = pd.concat((results, eval_df))
eval_df = evaluation('Ridge Classifier on Test', y1_Test, y1_Test_pred)
results = pd.concat((results, eval_df))

Ridge Classifier on Train
precision recall f1-score support
0 1.00 0.94 0.97 6995
1 0.94 1.00 0.97 6505

accuracy 0.97 0.97 0.97 13500
macro avg 0.97 0.97 0.97 13500
weighted avg 0.97 0.97 0.97 13500

Ridge Classifier on Test
precision recall f1-score support
0 1.00 0.94 0.97 2998
1 0.01 1.00 0.02 2


accuracy 0.51 0.97 0.94 3000
macro avg 1.00 0.94 0.97 3000
weighted avg 1.00 0.94 0.97 3000

In [64]: conf_mat(y1_Train, y1_Train_pred, 'Train', 'Ridge Clf')
Ridge Clf - Confusion Matrix on Train

Out[64]:
predicted_Non_Default predicted_Default
Non_Default 6556 439
Default 0 6505

In [65]: conf_mat(y1_Test, y1_Test_pred, 'Test', 'Ridge Clf')
Ridge Clf - Confusion Matrix on Test

Out[65]:
predicted_Non_Default predicted_Default
Non_Default 2804 194
Default 0 2
```



```
In [67]: results

Out[67]:
model Accuracy Precision Recall F1-score AUC
0 Random Forest Clf on Test 1 1 1 1 1
0 Random Forest Clf on Train 1 1 1 1 1
0 lightgbm on Train 0.98 0.98 1 0.98 0.98
0 lightgbm on Test 0.97 1 1 0.98 0.98
0 Logistic Regression on Train 0.94 0.95 1 0.94 0.94
0 Logistic Regression 0.88 1 1 0.01 0.94
0 LinearSVC on Train 0.84 0.86 0.72 0.81 0.84
0 LinearSVC on Test 0.95 1 1 0.03 0.97
0 Ridge Classifier on Train 0.97 0.97 1 0.97 0.97
0 Ridge Classifier on Test 0.94 1 1 0.02 0.97

In [ ]:
```